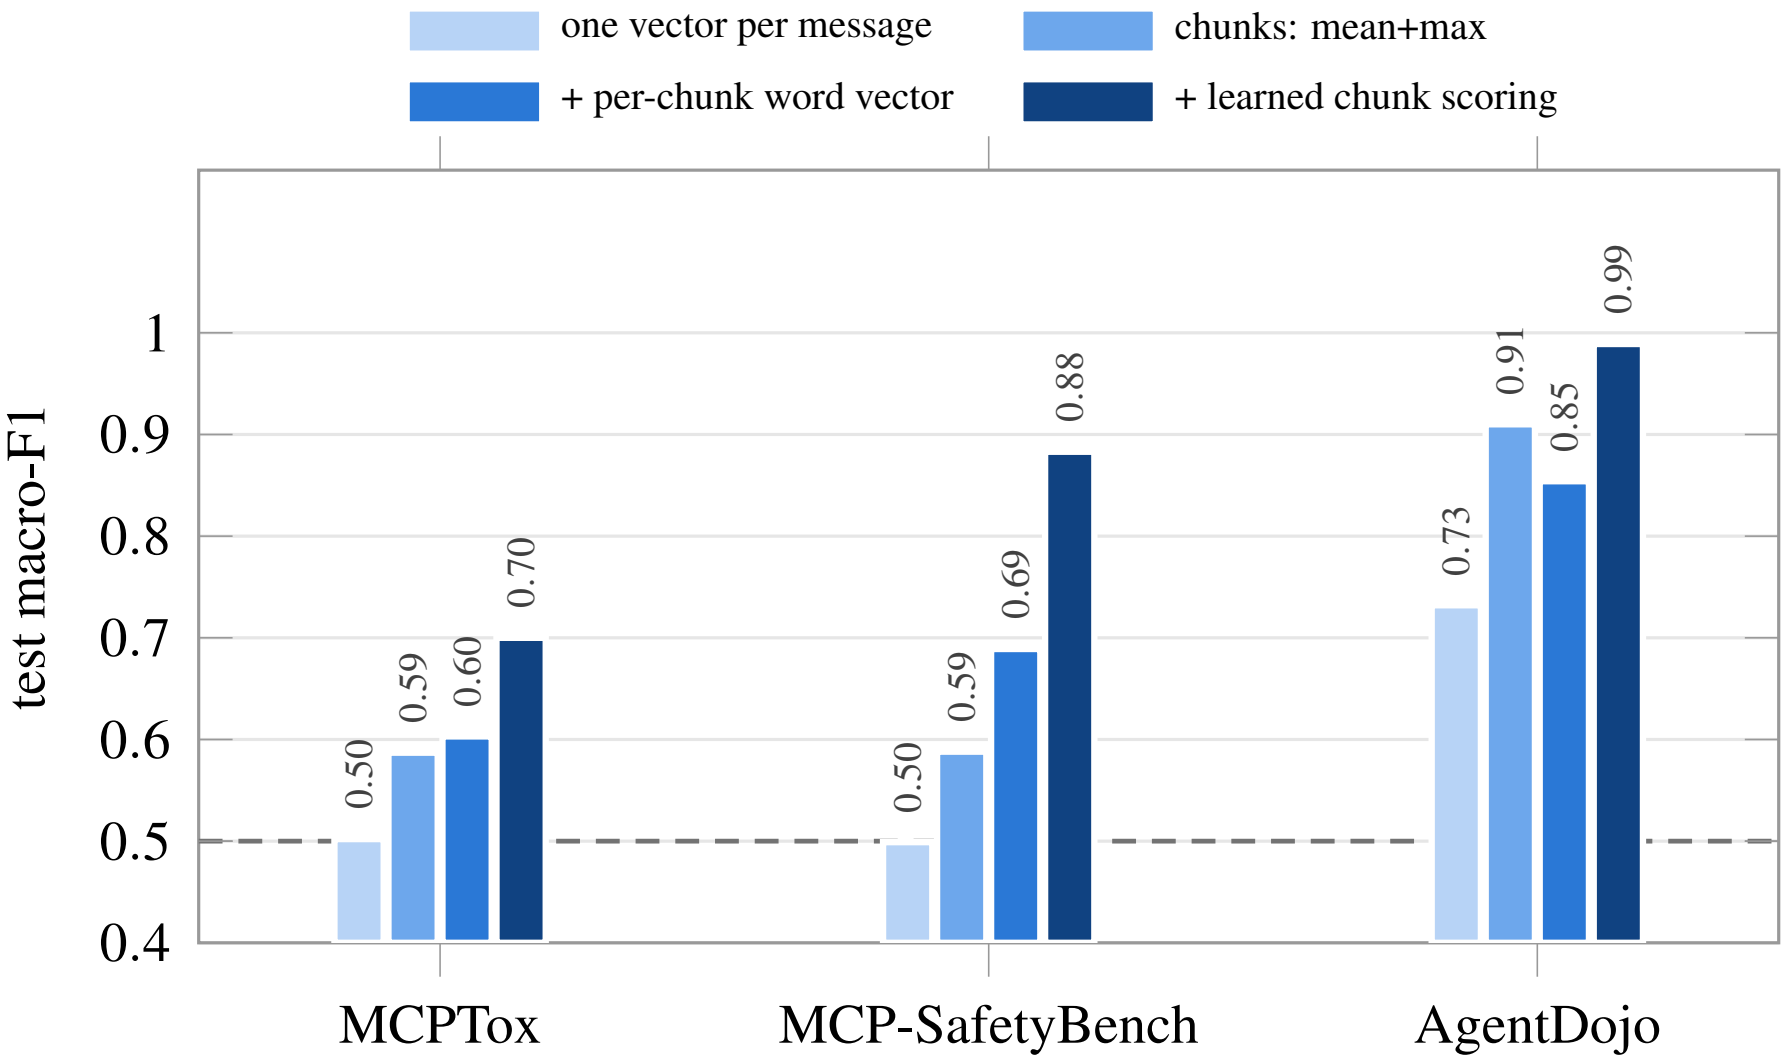
one vector per message
chunks: mean+max
+ per-chunk word vector
+ learned chunk scoring
test macro-F1
1
0.9
0.8
0.7
0.6
0.5
0.4
0.50
0.59
0.60
0.70
0.50
0.59
0.69
0.88
0.73
0.91
0.85
0.99
MCPTox
MCP-SafetyBench
AgentDojo